


# **Deep Learning Tutorial # 2 CNN using Keras**

CS 5102 Deep Learning

# Agenda

1. Train a model that would differentiate images of Cats from Images of Dogs.
2. Implement CNN using Keras
3. Project - composed of 3 parts, one week per part.

## Deep Learning Framework Power Scores 2018



# Dataset - Dogs Vs. Cats

1. Create an algorithm to distinguish dogs from cats.
2. Contains 25,000 images of dogs and cats.
3. Approach: Preprocessing followed by model training

# Dogs Vs. Cats

```
In [2]: import numpy as np
...: import pandas as pd
...: import cv2
...: import os
...: import matplotlib.pyplot as plt
```

```
In [3]: train = 'G:/Dogs vs. Cats/input/train/'
...: test = 'G:/Dogs vs. Cats/input/test1/'
```

# Dogs Vs. Cats


In [5]: *# Global Variables*

```
...: data = []
...: test_data = []
...: test_id = []
...: label = []
...: im_width = 64
...: im_height = 64
...: train_image_files = [ f for f in os.listdir(train)
...:                        if os.path.isfile(os.path.join(train,f)) ]
...: test_image_files = [ f for f in os.listdir(test)
...:                      if os.path.isfile(os.path.join(test,f)) ]
...:
```

# Dogs Vs. Cats

```
In [6]: image_file = str(train + train_image_files[5])  
...: img = cv2.imread(image_file)  
...: plt.imshow(img)  
...:
```


```
Out[6]: <matplotlib.image.AxesImage at 0x1fa89a8ce10>
```



# Dogs Vs. Cats

```
In [7]: image_file = str(train + train_image_files[22000])  
...: img = cv2.imread(image_file)  
...: plt.imshow(img)  
...:
```


```
Out[7]: <matplotlib.image.AxesImage at 0x1fa8996ae48>
```





# Dogs Vs. Cats

```
In [13]: image_file = str(train + train_image_files[24])  
....: img = cv2.imread(image_file, cv2.IMREAD_GRAYSCALE) # read a grayscale image  
....: plt.imshow(img, cmap='gray') # show a grayscale image  
Out[13]: <matplotlib.image.AxesImage at 0x143de8f7fd0>
```



# Dogs Vs. Cats

```
In [8]: def preprocessing(path):
...:     progress = 0
...:     image_files = [ f for f in os.listdir(path) if os.path.isfile(os.path.join(path,f)) ]
...:
...:     for file_name in image_files:
...:         image_file = str(path + file_name)
...:
...:         img = cv2.imread(image_file, cv2.IMREAD_GRAYSCALE)
...:         new_img = cv2.resize(img, (im_width, im_height))
...:
...:         if file_name[:3] == 'cat':
...:             label.append(0)
...:
...:         elif file_name[:3] == 'dog':
...:             label.append(1)
...:
...:         progress = progress + 1
...:         if progress % 1000 == 0:
...:             print('Progress: ' + str(progress) + ' Images Done')
...:
...:     print(len(data))
...:     print(len(label))
```

# Dogs Vs. Cats

```
In [3]: preprocessing(train)
```

```
Progress: 1000 Images Done  
Progress: 2000 Images Done  
Progress: 3000 Images Done  
Progress: 4000 Images Done  
Progress: 5000 Images Done  
Progress: 6000 Images Done  
Progress: 7000 Images Done  
Progress: 8000 Images Done  
Progress: 9000 Images Done  
Progress: 10000 Images Done  
Progress: 11000 Images Done  
Progress: 12000 Images Done  
Progress: 13000 Images Done  
Progress: 14000 Images Done  
Progress: 15000 Images Done  
Progress: 16000 Images Done  
Progress: 17000 Images Done  
Progress: 18000 Images Done  
Progress: 19000 Images Done  
Progress: 20000 Images Done  
Progress: 21000 Images Done  
Progress: 22000 Images Done  
Progress: 23000 Images Done  
Progress: 24000 Images Done  
Progress: 25000 Images Done  
25000  
25000
```

# CNN using Keras with Tensorflow Backend

```
In [4]: from keras import Sequential
...: from keras.layers import Dense,MaxPooling2D,Conv2D,Flatten,Dropout
...:
...: model = Sequential() # linear stack of layers.
...:
...: # Stride: stride length of the convolution
...: # Padding:
...: # A zero padding is used such that the output has the same length as the original input.
...: # One of "valid", "causal" or "same"
...: model.add(Conv2D(kernel_size=(3,3),filters=3,input_shape=(im_width, im_height, 1),
...:                    activation="relu",padding="valid"))
...: model.add(Conv2D(kernel_size=(3,3),filters=10,activation="relu",padding="same"))
...: model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
...:
...: model.add(Conv2D(kernel_size=(3,3),filters=3,activation="relu",padding="same"))
...: model.add(Conv2D(kernel_size=(5,5),filters=5,activation="relu",padding="same"))
...: model.add(MaxPooling2D(pool_size=(3,3),strides=(2,2)))
Using TensorFlow backend.
```

# CNN using Keras with Tensorflow Backend

```
In [5]: model.add(Conv2D(kernel_size=(2,2),strides=(2,2),filters=10))
...:
...: model.add(Flatten())
...: model.add(Dropout(0.2))
...: model.add(Dense(100,activation="sigmoid"))
...: model.add(Dense(1,activation="sigmoid"))
```

# CNN using Keras with Tensorflow Backend

```
In [6]: model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 62, 62, 3)	30
conv2d_2 (Conv2D)	(None, 62, 62, 10)	280
max_pooling2d_1 (MaxPooling2D)	(None, 31, 31, 10)	0
conv2d_3 (Conv2D)	(None, 31, 31, 3)	273
conv2d_4 (Conv2D)	(None, 31, 31, 5)	380
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 5)	0
conv2d_5 (Conv2D)	(None, 7, 7, 10)	210
flatten_1 (Flatten)	(None, 490)	0
dropout_1 (Dropout)	(None, 490)	0
dense_1 (Dense)	(None, 100)	49100
dense_2 (Dense)	(None, 1)	101
Total params: 50,374		
Trainable params: 50,374		
Non-trainable params: 0		



# CNN using Keras with Tensorflow Backend

```
In [7]: model.compile(optimizer="adadelta",loss="binary_crossentropy",metrics=["accuracy"])
```

```
In [8]: data = np.array(data)
...: print(data.shape)
(25000, 64, 64)
```

```
In [9]: data = data.reshape((data.shape)[0], (data.shape)[1], (data.shape)[2], 1)
...: print(data.shape)
(25000, 64, 64, 1)
```

```
In [10]: label = np.array(label)
...: print(label.shape)
(25000,)
```

# CNN using Keras with Tensorflow Backend

```
In [12]: model.fit(data, label, batch_size=128, epochs=2, verbose=1, validation_split=0.3)
```

```
Train on 17500 samples, validate on 7500 samples
```

```
Epoch 1/2
```

```
17500/17500 [=====] - 856s 49ms/step - loss: 0.6153 - acc: 0.7009 - val_loss: 1.2167 - val_acc: 0.0000e+00
```

```
Epoch 2/2
```

```
17500/17500 [=====] - 851s 49ms/step - loss: 0.5975 - acc: 0.7140 - val_loss: 1.0837 - val_acc: 0.0000e+00
```

```
Out[12]: <keras.callbacks.History at 0x143deba02b0>
```



# CNN using Keras with Tensorflow Backend

```
In [14]: # model.save("G:/Dogs vs. Cats/keras_model_attempt_1.h5")
```

```
In [15]: # model = load_model("G:/Dogs vs. Cats/keras_model_attempt_1.h5")
```

```
In [16]: def preprocessing_test(path):
...:     progress = 0
...:
...:     for file in test_image_files:
...:         image_file = str(path + file)
...:
...:         img = cv2.imread(image_file, cv2.IMREAD_GRAYSCALE)
...:         new_img = cv2.resize(img, (im_width, im_height))
...:
...:         test_id.append(file[:-4])
...:         test_data.append(new_img / 255)
...:
...:         progress = progress + 1
...:         if progress % 1000 == 0:
...:             print('Progress: ' + str(progress) + ' Images Done')
...:
...:     print(len(test_data))
...:     print(len(test_id))
...:
```

# CNN using Keras with Tensorflow Backend

```
In [17]: preprocessing_test(test)
```

```
Progress: 1000 Images Done  
Progress: 2000 Images Done  
Progress: 3000 Images Done  
Progress: 4000 Images Done  
Progress: 5000 Images Done  
Progress: 6000 Images Done  
Progress: 7000 Images Done  
Progress: 8000 Images Done  
Progress: 9000 Images Done  
Progress: 10000 Images Done  
Progress: 11000 Images Done  
Progress: 12000 Images Done  
12500  
12500
```

```
In [18]: test_data = np.array(test_data)
```

```
...: print(test_data.shape)
```

```
...:
```

```
(12500, 64, 64)
```

```
In [19]: test_data = test_data.reshape((test_data.shape)[0], (test_data.shape)[1], (test_data.shape)[2], 1)
```

```
...: print(test_data.shape)
```

```
...:
```

```
(12500, 64, 64, 1)
```

# CNN using Keras with Tensorflow Backend

```
In [20]: predicted_labels = model.predict(test_data)
```

```
In [21]: final_labels = []
...: for value in predicted_labels:
...:     if value > 0.5:
...:         final_labels.append(1)
...:     else:
...:         final_labels.append(0)
...:
```

```
In [22]: final_submission = pd.DataFrame({"id" : test_id })
```

```
In [23]: final_submission["label"] = final_labels
```

```
In [24]: final_submission.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12500 entries, 0 to 12499
Data columns (total 2 columns):
id      12500 non-null object
label   12500 non-null int64
dtypes: int64(1), object(1)
memory usage: 195.4+ KB
```

# CNN using Keras with Tensorflow Backend

```
In [25]: final_submission.head()
```

```
Out[25]:
```

	id	label
0	1	0
1	10	0
2	100	0
3	1000	0
4	10000	0

# Project One - MNIST Classifier

label = 5



label = 0



label = 4



label = 1



label = 9



label = 2



label = 1



label = 3



label = 1



label = 4



label = 3



label = 5



label = 3



label = 6



label = 1



label = 7



label = 2



label = 8



label = 6



label = 9



# Project One - MNIST Classifier


1. Composed of Three Parts
2. Each Part has a separate deadline
3. Dataset: MNIST Digit Dataset
4. The MNIST dataset comprises 60,000 training examples and 10,000 test examples of the handwritten digits 0–9, formatted as 28x28-pixel monochrome images.
5. Source: <http://yann.lecun.com/exdb/mnist/>
6. Alternate Source:

```
In [28]: from keras.datasets import mnist
...: (x_train, y_train), (x_test, y_test) = mnist.load_data()
...:
```

# Project One - MNIST Classifier

Part 1: Convolutional Neural Network using Keras with Tensorflow Backend.

Implement CNN in Keras following the LeNet Architecture from Yann LeCun.



Deadline: 1st November 2018

# Project One - MNIST Classifier

## Part 2: Visualization:

1. Plot error vs epochs.
2. Visualize the architecture and plot the weights.
3. Use Tensorboard.
4. Save results as an image.

Deadline: 8th November 2018

## Resources:

1. <http://fizzylogic.nl/2017/05/08/monitor-progress-of-your-keras-based-neural-network-using-tensorboard/>
2. <https://keras.io/callbacks/>
3. <https://www.youtube.com/watch?v=eBbEDRsCmv4>



# Project One - MNIST Classifier

## Part 3: Hyperparameter Tuning:

1. Common training problems: oscillating error, overfit.
2. Readjust the learning rate, weight initialization, and momentum.
3. Identify Overfit and Underfit.
4. Recreate the architecture with minimum layers while accuracy remains more than 90 %.
5. Recreate the architecture with maximum layers until the accuracy no longer improves significantly.
6. Write a report, more instructions will be sent shortly.

Deadline: 15th November 2018